
title: “The GEO Playbook: How to Get Your Brand Cited by AI” subtitle: “The Definitive Guide to Generative Engine Optimisation” version: “2.0” date: “February 2026”

The GEO Playbook ## How to Get Your Brand Cited by AI *The Definitive Guide to Generative Engine Optimisation* --- **February 2026 · v2.0**

“The best place to hide a dead body is page two of Google. In 2026, the best place to hide one is outside the AI’s answer.”

Table of Contents


1. The Shift: Why Search Will Never Be the Same
 2. What Is Generative Engine Optimisation (GEO)?
 3. How AI Engines Select Their Sources
 4. The Four Engines: A Platform-by-Platform Breakdown
 5. The 12 Pillars of GEO-Ready Content
 6. Before & After: Real Optimisation Examples
 7. The GEO Audit Checklist
 8. Measuring AI Visibility
 9. GEO vs. SEO: Allies, Not Rivals
 10. What’s Coming Next
-

1. The Shift: Why Search Will Never Be the Same

For twenty-five years, the game was simple: rank on Google, win the click, earn the customer. Entire industries — content marketing, link building, technical SEO — were built around a single behaviour: a person types a query, scans ten blue links, and chooses one.

That behaviour is dying.

ChatGPT surpassed 100 million users faster than any application in history. Google's AI Overviews now appear across billions of searches every month. Perplexity processes millions of research queries daily. Claude is the go-to reasoning engine for professionals making high-stakes decisions. And users aren't clicking through to websites — they're reading the AI's synthesised answer and moving on.



The implications are profound:

- **Traffic from traditional search is declining.** Multiple studies show click-through rates dropping 20–60% on queries where AI Overviews appear.
- **Brand discovery is shifting.** When a user asks “*What's the best project management tool for remote teams?*”, the AI doesn't return a list of links — it names specific brands. If yours isn't mentioned, you don't exist.

- **The citation is the new click.** Being cited as a source inside an AI-generated response is becoming more valuable than ranking #1 on a traditional SERP.

This isn't a prediction. It's already happening. And the brands that understand how to optimise for this new reality will dominate the next decade.

Welcome to Generative Engine Optimisation.

2. What Is Generative Engine Optimisation (GEO)?

Generative Engine Optimisation (GEO) is the practice of optimising your content, brand presence, and digital footprint so that AI-powered search engines — ChatGPT, Claude, Perplexity, Google AI Overviews, Gemini, and others — cite, reference, or recommend your brand in their generated responses.

The term was first formalised in a landmark 2023 research paper from Princeton, Georgia Tech, the Allen Institute, and IIT Delhi, later accepted at KDD 2024 (one of the most prestigious data science conferences in the world). The researchers demonstrated that targeted optimisation strategies could **boost content visibility in generative engine responses by up to 40%**.

GEO in One Sentence

SEO gets you ranked. GEO gets you cited.

Why “Optimisation” and Not “Marketing”?

GEO is deliberately framed as *optimisation* because, like SEO before it, it involves understanding how a system works and structuring your content to align with that system's selection criteria. It's not about gaming — it's about making your content *genuinely useful* to AI engines that are trying to give their users the best possible answer.

The Three Layers of GEO

Layer	What It Covers	Why It Matters
Content Layer	How your pages are written, structured, and formatted	AI engines need to parse, understand, and extract from your content
Authority Layer	Brand mentions, citations, reviews, and trust signals across the web	AI engines cross-reference multiple sources to validate claims
Technical Layer	Crawlability, schema markup, server-side rendering, and robots.txt configuration	AI crawlers must be able to access and interpret your content

3. How AI Engines Select Their Sources

Understanding *why* an AI engine cites one source over another is the foundation of GEO. While the exact algorithms are proprietary, extensive testing, published research, and reverse-engineering have revealed consistent patterns.

3.1 The Retrieval-Augmented Generation (RAG) Pipeline

Most modern AI search engines use a process called **Retrieval-Augmented Generation (RAG)**:

- 1. Query decomposition.** The user's prompt is broken into sub-queries. A question like "*What's the best CRM for small businesses in the UK?*" might become three separate searches: "top CRM software," "CRM for small businesses," and "UK CRM tools."
- 2. Source retrieval.** The engine searches its index (or the live web) for relevant sources. This step resembles traditional search — domain authority, relevance, and freshness all matter.
- 3. Relevance scoring.** Retrieved sources are scored for topical relevance, authority, recency, and specificity.
- 4. Synthesis and generation.** The language model reads the top-scoring sources and generates a response, weaving in information from multiple pages.
- 5. Citation assignment.** Some engines (Perplexity, Google AI Overviews) attach source citations to specific claims. Others (ChatGPT, Claude) may mention sources conversationally or provide them when asked.

3.2 The Seven Signals AI Engines Use to Select Sources

Based on the GEO research paper, industry testing, and our own analysis, here are the primary signals:

The Seven Signals AI Engines Use to Select Sources



Signal	Description	Impact
Topical Authority	Depth and breadth of coverage on a specific topic across your domain	★★★★★
Citation Density	How often your brand or domain is mentioned across external sources	★★★★★
Content Structure	Clear headings, direct answers, lists, and tables that LLMs can parse	★★★★★☆
Statistical Evidence	Inclusion of data points, percentages, and research-backed claims	★★★★★☆
Recency	How recently content was published or updated	★★★★★☆
E-E-A-T Signals	Demonstrated Experience, Expertise, Authoritativeness, and Trustworthiness	★★★★★☆
Quotability	Concise, well-phrased statements that can be extracted as standalone claims	★★★★☆☆

3.3 The “Consensus Effect”

AI engines don't just find information — they look for **consensus**. If multiple independent sources agree on a claim, the AI is far more likely to present it as fact. This is why brand mentions across diverse, authoritative sources matter enormously. A single brilliant blog post won't move the needle. Twenty mentions across industry publications, review sites, forums, and expert blogs will.

4. The Four Engines: A Platform-by-Platform Breakdown

Each major AI engine has distinct retrieval behaviours, data sources, and citation patterns. Optimising for all four requires understanding their differences.

4.1 ChatGPT (OpenAI)

Attribute	Detail
Data sources	Bing search index, GPTBot crawler, licensed publisher data, training data (pre-cutoff)
Citation style	Inline citations when browsing is active; conversational mentions from training data
Crawler	GPTBot (User-Agent: <code>GPTBot</code>)
Key ranking factors	Brand authority in training data, Bing ranking position, content depth

Optimisation priorities: - Ensure GPTBot is **not blocked** in your robots.txt - Build brand presence on platforms heavily represented in OpenAI's training data (Wikipedia, Reddit, Stack Overflow, major publications) - Optimise for Bing — ChatGPT's web browsing feature uses Bing as its search backbone - Include clear, quotable definitions and direct answers in your content

4.2 Claude (Anthropic)

Attribute	Detail
Data sources	Web search (when enabled), training data, uploaded documents

Attribute	Detail
Citation style	Conservative — cites sources when using web search; otherwise draws from training knowledge
Crawler	ClaudeBot (User-Agent: ClaudeBot)
Key ranking factors	Quality and depth of training data representation, accuracy, nuance

Optimisation priorities: - Allow ClaudeBot access in your robots.txt - Focus on long-form, nuanced, expert content — Claude is trained to value depth and accuracy - Ensure your brand appears in high-quality, well-structured publications and technical resources - Include balanced perspectives, caveats, and evidence-based reasoning — Claude's training emphasises truthfulness

4.3 Perplexity AI

Attribute	Detail
Data sources	Real-time web search (multiple search engines), its own index
Citation style	Numbered inline citations with direct source links — the most transparent of all engines
Crawler	PerplexityBot (User-Agent: PerplexityBot)
Key ranking factors	Topical relevance, content freshness, domain authority, structured data

Optimisation priorities: - Perplexity is the **most citation-heavy** engine — it cites 5–15 sources per response - Freshness matters enormously — Perplexity favours recently published or updated content - Structure content with clear H2/H3 headings that match natural-language questions - Include statistics, data points, and named sources — Perplexity preferentially surfaces data-rich content - Allow PerplexityBot in your robots.txt

4.4 Google AI Overviews (SGE)

Attribute	Detail
Data sources	Google's search index, Knowledge Graph, licensed content
Citation style	Source cards with thumbnails displayed alongside the AI Overview
Crawler	Googlebot (standard crawler)
Key ranking factors	Traditional Google ranking signals + E-E-A-T + structured data

Optimisation priorities: - AI Overviews predominantly cite pages already ranking in the **top 10 organic results** — SEO is a prerequisite - Schema markup (FAQ, HowTo, Article, Product) significantly increases the chance of citation - Optimise your Google Business Profile for local and commercial queries - AI Overviews pull heavily from content with clear, direct answers in the opening paragraph - Pages with original research, proprietary data, or expert quotes are cited more frequently

4.5 Cross-Platform Summary

AI Engine Comparison: Citation Behaviour

	ChatGPT	Claude	Perplexity	Google AI Overviews
Real-time Web Access	No	No	Yes	Yes
Inline Citations	Sometimes	Sometimes	Yes	Yes
Favours Fresh Content	No	No	Yes	Yes
Key Search Backbone	Proprietary Model	Proprietary Model	Multiple Search Engines	Google Search
Citation Frequency	Low	Low	High	High

Factor	ChatGPT	Claude	Perplexity	Google AI Overviews
Real-time web access	🔗 (via Bing)	🔗 (when enabled)	🔗 (primary)	🔗 (via Google index)

Factor	ChatGPT	Claude	Perplexity	Google AI Overviews
Inline citations	Sometimes	Sometimes	Always	Always
Favours fresh content	Moderate	Low	High	High
Respects robots.txt	Yes	Yes	Yes	Yes
Key search backbone	Bing	Various	Multiple	Google

5. The 12 Pillars of GEO-Ready Content

These are the actionable strategies that move the needle. Each pillar is grounded in published research, platform testing, and real-world results.

Pillar 1: Lead with a Direct, Quotable Answer

AI engines scan for concise, extractable answers — typically within the first 100 words of a page. Open every article with a clear, definitional statement.

Rule of thumb: *If an AI can't quote your opening paragraph as a standalone answer, rewrite it.*

Pillar 2: Structure for Machine Readability


Use semantic HTML and clear hierarchy:

- **H1** for the page title (one per page)
- **H2** for major sections
- **H3** for sub-topics within sections
- **Bulleted and numbered lists** for steps, features, and comparisons
- **Tables** for structured comparisons
- **Bold text** for key terms and definitions

AI engines parse structured content far more effectively than dense paragraphs.

Pillar 3: Embed Statistics and Data Points

GEO Optimisation Impact on AI Visibility



The GEO research paper found that content containing statistics and cited data achieved **30–40% higher visibility** in AI-generated responses compared to content without them. Every major claim should be supported by a number.

- “Our product significantly reduces onboarding time.”
- “Our product reduces onboarding time by 47%, based on a 2025 study of 1,200 enterprise users.”

Pillar 4: Include Expert Quotations

Content with attributed expert quotes performs measurably better in generative engines. The original research showed that adding quotations from credible sources improved visibility by up to **25%** in GEO benchmarks.

Name real people. Include their credentials. Make the quotes specific and substantive.

Pillar 5: Build Topical Authority Through Content Clusters

Don’t write one article — build a **content cluster**. A comprehensive hub page linking to 10–20 supporting articles on related sub-topics signals deep expertise to AI engines.

```
Hub: "Complete Guide to Remote Work Tools"
└─ "Best Video Conferencing Platforms 2026"
└─ "How to Choose a Project Management Tool"
└─ "Remote Team Communication: Slack vs Teams vs Discord"
└─ "Asynchronous Work: A Manager's Guide"
└─ (15 more supporting articles)
```

Pillar 6: Earn Brand Mentions Everywhere

Unlike traditional SEO, **unlinked brand mentions** appear to carry significant weight in GEO. AI engines don't need a hyperlink to register that your brand is associated with a topic — they just need to see your name mentioned in context.

Priority platforms for brand mentions: - Industry publications and blogs - Reddit threads and community discussions - YouTube video descriptions and transcripts - Podcast show notes and transcripts - Wikipedia (if your brand is notable enough) - User-generated review platforms (G2, Trustpilot, Capterra)

Pillar 7: Optimise for Conversational Queries

People ask AI engines questions differently than they type into Google. Queries are longer, more conversational, and often multi-part.

Traditional search: “*best CRM small business*” AI query: “*I run a 12-person marketing agency in Manchester. What CRM would you recommend and why?*”

Structure your content to answer the *intent behind the conversation*, not just the keyword.

Pillar 8: Publish Original Research and Proprietary Data

Nothing earns AI citations faster than being a **primary source**. If you can publish original surveys, benchmark reports, case studies with real numbers, or industry analysis that no one else has — you become the source that every AI must cite.

Pillar 9: Keep Content Ruthlessly Fresh

Perplexity and Google AI Overviews both heavily favour recent content. Implement a content refresh calendar:

- **Quarterly:** Update statistics, add new examples, revise recommendations
- **Bi-annually:** Restructure and expand based on new queries you're seeing
- **Annually:** Full rewrite with updated research and fresh expert contributions

Include visible “Last updated” dates — both AI crawlers and users value recency signals.

Pillar 10: Implement Comprehensive Schema Markup

Structured data helps AI engines understand *what* your content is, not just what it says.

Essential schema types for GEO: - `Article` — publication date, author, headline - `FAQPage` — question-and-answer pairs - `HowTo` — step-by-step processes - `Product` — specifications, pricing, reviews - `Organisation` — brand identity and contact details - `Person` — author expertise and credentials

Pillar 11: Ensure Technical Accessibility for AI Crawlers

AI crawlers differ from traditional search crawlers. Many struggle with:

- **Client-side JavaScript rendering** — Use server-side rendering (SSR) or static site generation (SSG) where possible
- **Paywalled or gated content** — AI crawlers typically cannot log in or bypass access restrictions
- **Heavily dynamic content** — Ensure critical information is in the initial HTML response

Robots.txt audit: Check that you’re not blocking AI crawlers. The major user agents to allow:

```
User-agent: GPTbot
Allow: /


User-agent: ClaudeBot
Allow: /


User-agent: PerplexityBot
Allow: /


User-agent: Google-Extended
Allow: /
```

Pillar 12: Write for Synthesis, Not Just Ranking

Traditional SEO content often buries the answer to pad word count. AI engines *extract and synthesise* — they don't reward length for its own sake.

Write every paragraph as if it might be the *only* paragraph the AI reads. Make each section self-contained, factual, and valuable in isolation.

6. Before & After: Real Optimisation Examples

Example 1: SaaS Product Page

Before (traditional marketing copy):

Welcome to CloudSync — the revolutionary platform that's transforming how teams collaborate in the modern workplace. Our innovative solution leverages cutting-edge technology to deliver an unparalleled experience. With CloudSync, your team will achieve new heights of productivity.

After (GEO-optimised):

CloudSync is a cloud-based team collaboration platform used by over 14,000 organisations across 80 countries. It combines real-time document editing, task management, and video conferencing in a single workspace. In a 2025 benchmark study by Forrester, teams using CloudSync reported a 34% reduction in context-switching and a 28% improvement in project delivery times compared to using separate tools.

Why it works: The optimised version contains a clear definition, specific numbers, a named research source, and factual claims an AI can confidently cite. The original is pure fluff — no AI would ever quote it.

Example 2: Blog Post Introduction

□ Before:

In today's fast-paced digital landscape, email marketing remains one of the most powerful tools in a marketer's arsenal. But are you really making the most of it? In this comprehensive guide, we'll explore everything you need to know about email marketing in 2026.

□ After:

Email marketing generates an average return of £36 for every £1 spent, making it the highest-ROI channel in digital marketing (DMA, 2025). This guide covers the five strategies that drove the best results in 2025–26: segmentation, behavioural triggers, interactive emails, AI-powered personalisation, and deliverability optimisation. Each section includes benchmarks, examples, and implementation steps.

Why it works: The optimised version opens with a cited statistic, previews specific content, and reads as an authoritative reference. An AI answering “*What is the ROI of email marketing?*” would cite this immediately.

Example 3: FAQ Section

▀ Before:

Q: How long does shipping take? A: Shipping times vary depending on your location and the shipping method you choose. We offer several options to meet your needs.

▀ After:

Q: How long does shipping take? A: Standard UK delivery takes 3–5 working days. Express delivery arrives within 1–2 working days. International shipping to Europe takes 5–10 working days and worldwide 10–21 working days. All orders include tracking. Orders placed before 2:00 PM GMT on weekdays are despatched the same day.

Why it works: Specificity. AI engines prioritise content that answers questions *completely* in a single passage. Vague answers are never cited.

7. The GEO Audit Checklist

Use this checklist to audit any page on your site for GEO readiness.

Content Quality

- [] Opening paragraph contains a clear, quotable definition or answer
- [] Key claims are supported with statistics, percentages, or named data sources
- [] Expert quotes with attributed names and credentials are included
- [] Content answers the query directly — no filler before the answer
- [] Each section is self-contained and extractable
- [] Content is written in British or American English consistently (not mixed)

- [] “Last updated” date is visible and accurate

Content Structure

- [] Proper heading hierarchy (H1 → H2 → H3, no skipped levels)
- [] Bulleted or numbered lists used for steps, features, and comparisons
- [] Tables used for structured comparisons where appropriate
- [] Key terms and definitions are bold or highlighted
- [] FAQ section included with question-format headings

Technical Foundation

- [] AI crawlers (GPTBot, ClaudeBot, PerplexityBot, Google-Extended) are allowed in robots.txt
- [] Page uses server-side rendering or pre-rendering
- [] Schema markup implemented (Article, FAQPage, HowTo, Product, Organisation as relevant)
- [] Page loads without requiring JavaScript for critical content
- [] Core Web Vitals passing (Page Speed, LCP, CLS, INP)
- [] HTTPS enabled
- [] Mobile responsive

Authority & Distribution

- [] Brand is mentioned on at least 10 external authoritative sources for the target topic
- [] Content is referenced or linked from industry publications
- [] Brand has presence on UGC platforms (Reddit, Quora, YouTube, G2)
- [] Google Business Profile is complete and optimised (for local/commercial queries)
- [] Wikipedia article exists or brand is mentioned in relevant Wikipedia articles
- [] Social proof (reviews, testimonials, case studies) is publicly accessible

Freshness

- [] Content has been updated within the last 90 days
- [] Statistics and data points reference the current or previous year

- [] No broken links or references to discontinued products/services
 - [] Content refresh schedule is documented and followed
-

8. Measuring AI Visibility

You can't improve what you can't measure. Here's how to track your GEO performance.

Manual Monitoring

The prompt test: Regularly ask each AI engine questions that your brand should appear in. Track:

- Whether your brand is mentioned
- Whether your website is cited as a source
- What position your brand appears in (first mentioned? third?)
- What competitors are cited instead
- How the AI describes your brand (accurate? favourable? outdated?)

Create a tracking spreadsheet:

Prompt	Engine	Brand Mentioned?	Position	Competitors Cited	Date
“Best CRM for agencies UK”	ChatGPT	☐ Yes	2nd	HubSpot, Pipedrive	03/02/26
“Best CRM for agencies UK”	Perplexity	☐ No	—	Salesforce, Zoho	03/02/26
“Best CRM for agencies UK”	Google AIO	☐ Yes	1st	HubSpot	03/02/26

Automated Tools

Several platforms now offer AI visibility tracking:

- **Semrush AI Visibility Toolkit** — Tracks brand mentions, sentiment, and share of voice across AI engines
- **Profound** — Monitors brand presence in LLM responses
- **Scrunch AI** — Analyses AI citation patterns and competitor visibility
- **Otterly.AI** — Tracks brand appearances in AI-generated answers
- **Peec AI** — Monitors AI search results for brand mentions

Key Metrics to Track

Metric	Definition	Target
AI Share of Voice	% of relevant prompts where your brand is cited vs. competitors	> 30%
Citation Frequency	How often your domain appears as a cited source	Increasing MoM
Brand Sentiment	Whether AI descriptions of your brand are positive, neutral, or negative	Positive
Source Accuracy	Whether AI engines present accurate information about your brand	> 95%
Prompt Coverage	% of target prompts where your brand appears at all	> 50%

9. GEO vs. SEO: Allies, Not Rivals

A common misconception is that GEO replaces SEO. It doesn't. They are deeply complementary.

What SEO Does for GEO

- **Ranking in Google's top 10 is a prerequisite** for being cited in AI Overviews — if you don't rank, you won't be cited
- **Backlinks build the domain authority** that AI engines use as a trust signal during source retrieval

- **Keyword research identifies the queries** that inform your GEO content strategy
- **Technical SEO ensures crawlability** for both traditional and AI crawlers

What GEO Adds Beyond SEO

- **Unlinked mention strategy** — building brand presence where links don't matter
- **Cross-platform optimisation** — ensuring visibility in ChatGPT, Claude, and Perplexity, not just Google
- **Synthesis-first writing** — structuring content for extraction rather than just ranking
- **AI crawler management** — specifically managing access for GPTBot, ClaudeBot, and PerplexityBot
- **Prompt-based strategy** — optimising for conversational queries, not just keyword strings

The Unified Approach

The most effective strategy is to treat GEO as an **extension** of your existing SEO programme. Every SEO improvement you make — better content, more backlinks, stronger technical foundation — also improves your GEO performance. The GEO-specific additions (structured answers, statistical density, expert quotes, unlinked mentions, AI crawler access) layer on top.

10. What's Coming Next

GEO is in its infancy. Here's what we expect to see in the next 12–24 months:

1. **AI engine advertising.** Paid placements within AI-generated responses are coming. Perplexity has already begun testing sponsored results. Google will follow.
2. **Standardised AI analytics.** Just as Google Analytics became the standard for web traffic, tools for measuring AI citation performance will mature rapidly.
3. **AI-specific content formats.** We expect new structured data standards designed specifically for AI consumption — beyond current schema.org markup.
4. **Personalised AI results.** As AI engines learn individual user preferences, brand visibility will become more nuanced. Your content may be shown to some users and hidden from others based on their history.

- 5. Voice and multimodal search.** As AI assistants become the primary interface (through smart speakers, in-car systems, AR glasses), GEO will expand to optimise for voice synthesis and visual responses.
 - 6. Regulatory scrutiny.** Governments are beginning to ask whether AI engines have an obligation to fairly represent all sources, not just the ones they prefer. This could fundamentally reshape how citations work.
-

Final Thoughts

The transition from traditional search to AI-powered discovery is not a distant possibility — it is happening now. Every month, more users turn to ChatGPT, Claude, Perplexity, and Google AI Overviews for the answers that used to require clicking through ten blue links.

The brands that win in this new landscape will be those that understand a simple truth: **AI engines don't rank pages. They cite sources.** And to be cited, your content must be authoritative, structured, specific, evidence-based, and present across the web.

GEO is not a replacement for good marketing. It's the evolution of it. The fundamentals — genuine expertise, valuable content, trustworthy reputation — matter more than ever. GEO simply gives you the framework to ensure those fundamentals are visible to the machines that are increasingly deciding which brands get recommended and which get ignored.

Start with the checklist. Audit your top ten pages. Fix the gaps. Monitor your AI visibility. Iterate.

The first-mover advantage in GEO is real, and the window is open.

— End of Playbook — **The GEO Playbook** · v2.0 · February 2026 *© 2026. All rights reserved.*