

DIGITAL PRODUCT

The GEO Playbook

How to Get Your Brand Cited by AI

by Rook 

1.0 · February 2026

"The best place to hide a dead body is page two of Google. In 2026, the best place to hide one is outside the AI's answer."

Table of Contents

- 1 The Shift: Why Search Will Never Be the Same
- 2 What Is Generative Engine Optimisation (GEO)?
- 3 How AI Engines Select Their Sources
- 4 The Four Engines: A Platform-by-Platform Overview
- 5 Platform Deep Dives: Engine-by-Engine Optimisation
- 6 The 12 Pillars of GEO-Ready Content
- 7 Before & After: Real Optimisation Examples
- 8 The Content Audit Framework
- 9 The GEO Audit Checklist
- 10 Implementation Roadmap: Your First 30 Days
- 11 Measuring AI Visibility
- 12 GEO vs. SEO: Allies, Not Rivals
- 13 What's Coming Next
- 14 Appendix: Resources, Tools, and Further Reading

1. The Shift: Why Search Will Never Be the Same

For twenty-five years, the game was simple: rank on Google, win the click, earn the customer. Entire industries, from content marketing to link building to technical SEO, were built around a single behaviour: a person types a query, scans ten blue links, and chooses one.

That behaviour is dying.

ChatGPT surpassed 100 million users faster than any application in history, reaching that milestone in just two months. Google's AI Overviews now appear across billions of searches every month. Perplexity processes millions of research queries daily. Claude is the go-to reasoning engine for professionals making high-stakes decisions. And users are not clicking through to websites; they are reading the AI's synthesised answer and moving on.

► The Numbers That Should Worry Every Marketer

The data is unambiguous. Here is what published research and industry reporting tells us about the shift already underway:

METRIC	FINDING	SOURCE
Click-through rate decline	20-60% drop on queries where AI Overviews appear	Multiple industry studies, 2024-25
ChatGPT monthly active users	300M+ as of late 2025	OpenAI
Google AI Overviews reach	Appearing on billions of monthly searches	Google I/O 2025
Perplexity monthly queries	500M+ and growing rapidly	Perplexity, Q4 2025
Zero-click searches	Over 65% of Google searches now end without a click	SparkToro / Datos, 2025
AI search adoption	79% of consumers use AI search weekly	Botify State of AI Search, 2025

These are not projections. These are measurements of what is already happening.

► The Brand Discovery Problem

When a user asks "What's the best project management tool for remote teams?", the AI does not return a list of links. It names specific brands. If yours is not mentioned, you do not exist in that moment.

Consider this scenario. A marketing director in Birmingham needs a new CRM. Five years ago, she would have typed "best CRM for agencies" into Google, clicked three or four results, and compared features. Today, she opens ChatGPT and types: "I run a 40-person marketing agency. We need a CRM that integrates with HubSpot, handles complex deal pipelines, and costs under £80 per seat. What would you recommend?"

ChatGPT names three tools. If your CRM is not one of them, you never had a chance.

This is the new reality:

- **The citation is the new click.** Being cited as a source inside an AI-generated response is becoming more valuable than ranking #1 on a traditional SERP.
- **Brand discovery is conversational.** Users are asking AI for recommendations the way they would ask a trusted colleague.
- **Visibility is binary.** In a list of ten blue links, being #7 still gets some traffic. In an AI response that names three brands, being #4 means being invisible.

► Case Study: The SaaS Company That Disappeared

A mid-market project management platform (anonymised at their request) tracked their inbound leads from January to September 2025. Despite maintaining stable Google rankings (positions 3-6 for their primary keywords), they observed a 31% decline in organic traffic from those same keywords. When they investigated, they discovered that Google AI Overviews had been activated for 68% of their target queries, and their brand was cited in only 12% of those overviews. Meanwhile, two competitors, one of whom ranked lower in organic results, were cited in over 70% of the same AI Overviews because those competitors had invested heavily in structured content, expert commentary, and third-party brand mentions.

The lesson: organic rankings alone no longer guarantee traffic. AI citation is the new gatekeeper.

► What This Means for You

This is not a prediction. It is already happening. And the brands that understand how to optimise for this new reality will dominate the next decade. The good news is that the playbook for doing so is clear, actionable, and, for now, a genuine competitive advantage. Most businesses have not even begun thinking about this.

Welcome to Generative Engine Optimisation.

2. What Is Generative Engine Optimisation (GEO)?

Generative Engine Optimisation (GEO) is the practice of optimising your content, brand presence, and digital footprint so that AI-powered search engines, including ChatGPT, Claude, Perplexity, Google AI Overviews, Gemini, and others, cite, reference, or recommend your brand in their generated responses.

The term was first formalised in a landmark 2023 research paper from Princeton, Georgia Tech, the Allen Institute, and IIT Delhi, later accepted at KDD 2024, one of the most prestigious data science conferences in the world. The researchers demonstrated that targeted optimisation strategies could **boost content visibility in generative engine responses by up to 40%**.

► GEO in One Sentence

SEO gets you ranked. GEO gets you cited.

► Why "Optimisation" and Not "Marketing"?

GEO is deliberately framed as optimisation because, like SEO before it, it involves understanding how a system works and structuring your content to align with that system's selection criteria. It is not about gaming; it is about making your content genuinely useful to AI engines that are trying to give their users the best possible answer.

► The Three Layers of GEO

LAYER	WHAT IT COVERS	WHY IT MATTERS	KEY ACTIONS
Content Layer	How your pages are written, structured, and formatted	AI engines need to parse, understand, and extract from your content	Restructure existing pages, add statistics, include quotable definitions
Authority Layer	Brand mentions, citations, reviews, and trust signals across the web	AI engines cross-reference multiple sources to validate claims	Build mentions on review sites, earn press coverage, grow forum presence
Technical Layer	Crawlability, schema markup, server-side rendering, and robots.txt configuration	AI crawlers must be able to access and interpret your content	Audit robots.txt, implement schema, ensure SSR

► The GEO Maturity Model

Most organisations fall into one of four stages. Understanding where you are helps you prioritise what to do first.

STAGE	NAME	DESCRIPTION	TYPICAL AI VISIBILITY
1	Invisible	No GEO consideration; content written for humans only or for traditional SEO	0-5% of target prompts
2	Incidental	Some AI citations happen by accident, usually because of strong existing SEO	5-20% of target prompts
3	Intentional	Active GEO optimisation underway; content restructured, schema added, mentions growing	20-50% of target prompts
4	Dominant	Comprehensive GEO programme; cited consistently across engines for target topics	50%+ of target prompts

Most businesses reading this playbook are at Stage 1 or 2. The goal is to move to Stage 3 within 30 days and Stage 4 within 90 days.

► The Economics of GEO

Why does this matter commercially? Consider the maths:

SCENARIO	TRADITIONAL SEO	GEO-OPTIMISED
Monthly searches for target query	10,000	10,000
Traffic reaching AI answer (no click)	0% (users clicked)	60% (users read AI answer)
Users who see your brand in AI answer	N/A	4,200 (70% citation rate × 60%)
Users who click to your site from AI	N/A	630 (15% click-through from citation)
Users who click from organic result	2,500 (25% CTR for position 3)	1,000 (10% CTR, reduced by AI answer)
Total brand impressions	2,500	5,200
Total site visits	2,500	1,630

Note: total site visits may decrease, but total brand impressions increase dramatically. The user who reads the AI's recommendation of your product and then visits your site directly two days later is not captured by click attribution, yet the sale is real. GEO is as much a brand strategy as a traffic strategy.

3. How AI Engines Select Their Sources

Understanding why an AI engine cites one source over another is the foundation of GEO. While the exact algorithms are proprietary, extensive testing, published research, and reverse-engineering have revealed consistent patterns.

► 3.1 The Retrieval-Augmented Generation (RAG) Pipeline

Most modern AI search engines use a process called **Retrieval-Augmented Generation (RAG)**. Understanding this pipeline is essential because each step represents an opportunity to optimise.

Step-by-step breakdown:

- 1 **Query decomposition.** The user's prompt is broken into sub-queries. A question like "What's the best CRM for small businesses in the UK?" might become three separate searches: "top CRM software," "CRM for small businesses," and "UK CRM tools."
- 2 **Source retrieval.** The engine searches its index (or the live web) for relevant sources. This step resembles traditional search: domain authority, relevance, and freshness all matter.
- 3 **Relevance scoring.** Retrieved sources are scored for topical relevance, authority, recency, and specificity. Typically, the engine retrieves 20-50 candidate sources and narrows to the top 5-10.
- 4 **Synthesis and generation.** The language model reads the top-scoring sources and generates a response, weaving in information from multiple pages. This is where the AI decides what to say and who to credit.
- 5 **Citation assignment.** Some engines (Perplexity, Google AI Overviews) attach source citations to specific claims. Others (ChatGPT, Claude) may mention sources conversationally or provide them when asked.

Optimisation opportunities at each stage:

RAG STAGE	WHAT HAPPENS	YOUR OPTIMISATION LEVER
Query decomposition	Prompt split into sub-queries	Ensure content addresses specific sub-topics, not just broad themes
Source retrieval	Engine searches for relevant pages	SEO fundamentals: ranking, domain authority, indexing, crawlability
Relevance scoring	Sources scored and filtered	Content specificity, recency, statistical density, structural clarity
Synthesis	AI reads and combines sources	Quotable passages, clear definitions, extractable facts

RAG STAGE	WHAT HAPPENS	YOUR OPTIMISATION LEVER
Citation	AI attributes claims to sources	Brand name proximity to key claims, clear authorship, schema markup

► 3.2 The Seven Signals AI Engines Use to Select Sources

Based on the GEO research paper, industry testing, and extensive analysis, here are the primary signals that determine whether your content gets cited:

SIGNAL	DESCRIPTION	IMPACT	HOW TO OPTIMISE
Topical Authority	Depth and breadth of coverage on a specific topic across your domain	★★★★★	Build content clusters with 10-20 articles per core topic
Citation Density	How often your brand or domain is mentioned across external sources	★★★★★	Earn brand mentions on 20+ external sites per key topic
Content Structure	Clear headings, direct answers, lists, and tables that LLMs can parse	★★★★☆	Use semantic HTML, tables, and lists; lead with direct answers
Statistical Evidence	Inclusion of data points, percentages, and research-backed claims	★★★★☆	Add 3-5 statistics per page, all with named sources
Recency	How recently content was published or updated	★★★★☆	Update key pages quarterly; show visible "Last updated" dates
E-E-A-T Signals	Demonstrated Experience, Expertise, Authoritativeness, and Trustworthiness	★★★★☆	Include author bios, expert quotes, credentials, and real case studies
Quotability	Concise, well-phrased statements that can be extracted as standalone claims	★★★☆☆	Write "tweetable" opening sentences for every section

► 3.3 The "Consensus Effect"

AI engines do not just find information; they look for **consensus**. If multiple independent sources agree on a claim, the AI is far more likely to present it as fact. This is why brand mentions across diverse, authoritative sources matter enormously.

A single brilliant blog post will not move the needle. Twenty mentions across industry publications, review sites, forums, and expert blogs will.

How the consensus effect works in practice:

Imagine a user asks: "What is the best email marketing platform for e-commerce?"

The AI retrieves 30 sources. It finds that: - 18 sources mention Klaviyo as a top recommendation - 12 sources mention Mailchimp - 7 sources mention Omnisend - 2 sources mention your platform, "MailFlow"

The AI will almost certainly name Klaviyo first, Mailchimp second, and possibly Omnisend third. MailFlow, despite potentially having a superior product, does not get mentioned because it lacks the consensus signal.

The lesson: GEO is not just about optimising your own website. It is about ensuring your brand appears across the entire ecosystem of sources that AI engines consult.

► 3.4 Source Diversity and Trust Tiers

AI engines weight sources differently based on their perceived trust level. Through extensive testing, we can categorise sources into tiers:

TRUST TIER	SOURCE TYPE	EXAMPLES	WEIGHT
Tier 1: Institutional	Academic papers, government sites, established encyclopaedias	.gov, .edu, Wikipedia, PubMed	Highest
Tier 2: Editorial	Major publications with editorial oversight	Forbes, TechCrunch, BBC, industry journals	High
Tier 3: Expert	Recognised expert blogs, industry analysts, professional associations	Gartner, HBR blogs, professional body sites	High
Tier 4: Community	User-generated content with voting/curation mechanisms	Reddit, Stack Overflow, Quora, G2 reviews	Medium-High

TRUST TIER	SOURCE TYPE	EXAMPLES	WEIGHT
Tier 5: Brand	Company websites, brand blogs, product pages	Your own site	Medium
Tier 6: General	Generic blogs, article directories, low-authority sites	Content farms, thin affiliate sites	Low

Key insight: Your own website (Tier 5) carries less weight than external mentions (Tiers 1-4). This is why an external mention strategy is at least as important as on-site optimisation.

4. The Four Engines: A Platform-by-Platform Overview

Each major AI engine has distinct retrieval behaviours, data sources, and citation patterns. Optimising for all four requires understanding their differences.

► 4.1 ChatGPT (OpenAI)

ATTRIBUTE	DETAIL
Data sources	Bing search index, GPTBot crawler, licensed publisher data, training data (pre-cutoff)
Citation style	Inline citations when browsing is active; conversational mentions from training data
Crawler	GPTBot (User-Agent: GPTBot)
Key ranking factors	Brand authority in training data, Bing ranking position, content depth
Monthly active users	300M+ (as of late 2025)
Primary use case	General questions, recommendations, research, creative tasks

► 4.2 Claude (Anthropic)

ATTRIBUTE	DETAIL
Data sources	Web search (when enabled), training data, uploaded documents
Citation style	Conservative: cites sources when using web search; otherwise draws from training knowledge
Crawler	ClaudeBot (User-Agent: <code>ClaudeBot</code>)
Key ranking factors	Quality and depth of training data representation, accuracy, nuance
Primary use case	Professional analysis, complex reasoning, technical research, document review

► 4.3 Perplexity AI

ATTRIBUTE	DETAIL
Data sources	Real-time web search (multiple search engines), its own index
Citation style	Numbered inline citations with direct source links: the most transparent of all engines
Crawler	PerplexityBot (User-Agent: <code>PerplexityBot</code>)
Key ranking factors	Topical relevance, content freshness, domain authority, structured data
Monthly queries	500M+ (as of Q4 2025)
Primary use case	Fact-checking, research, current events, comparison shopping

► 4.4 Google AI Overviews (SGE)

ATTRIBUTE	DETAIL
Data sources	Google's search index, Knowledge Graph, licensed content
Citation style	Source cards with thumbnails displayed alongside the AI Overview
Crawler	Googlebot (standard crawler)
Key ranking factors	Traditional Google ranking signals + E-E-A-T + structured data
Reach	Billions of monthly searches
Primary use case	Quick answers, local queries, product comparisons, definitions

► 4.5 Cross-Platform Comparison Matrix

FACTOR	CHATGPT	CLAUDE	PERPLEXITY	GOOGLE AI OVERVIEWS
Real-time web access	■ (via Bing)	■ (when enabled)	■ (primary)	■ (via Google index)
Inline citations	Sometimes	Sometimes	Always	Always
Favours fresh content	Moderate	Low	High	High
Respects robots.txt	Yes	Yes	Yes	Yes
Key search backbone	Bing	Various	Multiple	Google
Number of sources cited	3-6	2-5	5-15	3-8
Training data influence	High	High	Low	Low

Factor	ChatGPT	Claude	Perplexity	Google AI Overviews
User base size	Largest	Growing	Growing fast	Largest (via Google Search)
Best for B2B	★★★★★	★★★★★	★★★★★	★★★★★
Best for B2C	★★★★★	★★★★★	★★★★★	★★★★★
Best for local	★★★★★	★★★★★	★★★★★	★★★★★

5. Platform Deep Dives: Engine-by-Engine Optimisation

This section provides specific, actionable optimisation tactics for each AI engine. While many GEO principles apply universally, each platform has unique characteristics that reward tailored strategies.

▶ 5.1 ChatGPT Deep Dive

Understanding ChatGPT's source selection:

ChatGPT draws from two distinct pools: its training data (for queries where browsing is not triggered) and live web search via Bing (for queries requiring current information). This dual-source nature means you need to optimise for both.

Training data optimisation:

ChatGPT's training data includes a vast corpus of web content, books, and licensed material. Brands that appeared frequently in high-quality content before the training cutoff have an inherent advantage. While you cannot retroactively change training data, you can:

- Ensure your brand has a strong, accurate Wikipedia presence
- Maintain an active, high-quality presence on platforms heavily represented in training data (Reddit, Stack Overflow, major news outlets)
- Publish on platforms with data licensing agreements with OpenAI (e.g., Associated Press, specific publishers)

Live search optimisation (Bing-dependent):

When ChatGPT browses the web, it uses Bing as its search backbone. This means Bing SEO directly influences ChatGPT citations.

BING OPTIMISATION TACTIC	WHY IT MATTERS FOR CHATGPT
Submit your sitemap to Bing Webmaster Tools	Ensures full indexing in Bing, which feeds ChatGPT
Optimise for Bing's ranking factors (social signals carry more weight than on Google)	Higher Bing rank = more likely to be retrieved by ChatGPT
Ensure fast page load times	Bing weights page speed heavily in its ranking algorithm
Use clear meta descriptions	ChatGPT often surfaces meta descriptions as part of its synthesis
Build links from .edu and .gov domains	Bing places a higher premium on institutional links than Google does

ChatGPT-specific content tactics:

- 1 Write "definition-first" content.** ChatGPT frequently quotes the first sentence of a page when defining a concept. Ensure your opening sentence is a clean, authoritative definition.
- 2 Create comprehensive comparison pages.** ChatGPT excels at comparative recommendations ("X vs Y" or "Best tools for Z"). Build thorough, balanced comparison content that the model can synthesise.
- 3 Include brand + category associations.** ChatGPT often relies on co-occurrence patterns. If your brand is consistently mentioned alongside your target category (e.g., "MailFlow, the email marketing platform for e-commerce"), the model learns that association.
- 4 Leverage Reddit strategically.** Reddit is one of the most heavily represented platforms in ChatGPT's training data. Genuine participation in relevant subreddits, including helpful answers that mention your brand naturally, can influence training-data-based recommendations.
- 5 Robots.txt configuration:**

User-agent: GPTBot
Allow: /

Testing protocol for ChatGPT:

Run these prompts monthly and track results:

TEST PROMPT	WHAT YOU ARE MEASURING
"What is [your product category]?"	Whether your brand appears in the definition
"What are the best [product category] tools?"	Whether you are in the recommendation list
"Compare [your brand] vs [competitor]"	How accurately the AI describes your product
"[Your brand] review"	What the AI says about your reputation
"I need a [product] that does [specific feature]. What do you recommend?"	Whether you appear for feature-specific queries

► 5.2 Claude Deep Dive

Understanding Claude's source selection:

Claude is built by Anthropic with a strong emphasis on accuracy, nuance, and safety. This makes Claude uniquely valuable for professional and B2B contexts: users trust Claude for high-stakes research, technical analysis, and balanced recommendations. Claude's web search, when enabled, tends to retrieve fewer but higher-quality sources.

What makes Claude different:

CHARACTERISTIC	IMPLICATION FOR GEO
Emphasis on accuracy and truthfulness	Content with hedging, caveats, and balanced perspectives performs better
Trained on high-quality, long-form content	In-depth articles outperform short-form listicles
Conservative citation style	Getting cited by Claude means your content passed a high quality bar
Strong at technical and professional topics	B2B and technical content gets disproportionate representation
Prefers primary sources	Original research and first-party data are strongly favoured

Claude-specific content tactics:

- 1 **Write with intellectual honesty.** Claude's training emphasises nuance. Content that acknowledges trade-offs, limitations, and alternative perspectives is more likely to be cited than content that makes absolute, unqualified claims.
- 2 **"Our platform is the best solution for every business."**
- 3 **"Our platform is particularly well-suited for teams of 20-100 with complex workflows, though businesses with simpler needs may find lighter-weight alternatives more cost-effective."**
- 4 **Include methodology sections.** When presenting data or research, explain how the data was collected. Claude values transparency and is more likely to cite sources that show their working.
- 5 **Build technical depth.** Claude users tend to ask more technical, specific questions than ChatGPT users. Ensure your content goes deep enough to answer questions like: "What's the architecture behind [your product]?" or "How does [your product] handle [specific edge case]?"
- 6 **Publish on platforms Claude respects.** Focus on:
- 7 Academic and research platforms (arXiv, Google Scholar)
- 8 Technical blogs and documentation
- 9 Industry journals with peer review or editorial standards
- 10 Professional association publications

11 Robots.txt configuration:

```
User-agent: ClaudeBot  
Allow: /
```

Claude user persona:

Understanding who uses Claude helps you tailor content for AI citation:

SEGMENT	TYPICAL QUERY	CONTENT NEEDED
Software developers	"What's the best approach to [technical problem]?"	Technical documentation, architecture guides, code examples
Business analysts	"Compare [solution A] vs [solution B] for [specific use case]"	Detailed comparison pages with real metrics
Researchers	"What does the evidence say about [topic]?"	Literature reviews, original research, meta-analyses
Consultants	"Create a framework for [business challenge]"	Frameworks, templates, methodology explanations
Legal/compliance	"What are the requirements for [regulation]?"	Regulatory guides, compliance checklists

► 5.3 Perplexity Deep Dive

Understanding Perplexity's source selection:

Perplexity is the most citation-transparent AI engine. Every response includes numbered inline citations linking directly to source URLs. This makes it the most measurable platform for SEO, and the one where optimisation has the most immediate, visible impact.

Perplexity's retrieval characteristics:

BEHAVIOUR	DETAIL
Sources per response	5-15 (significantly more than other engines)
Real-time indexing	Perplexity crawls the live web for every query; freshness is critical
Source diversity	Tends to cite a mix of authoritative sites and niche experts
Content format preference	Strong preference for structured, data-rich content
Update sensitivity	Content updated within the last 30 days has a measurable citation advantage

Perplexity-specific content tactics:

- 1 **Publish and update frequently.** Perplexity's real-time crawling means that a page updated yesterday will often outperform a higher-authority page updated six months ago. Implement a weekly content refresh schedule for your highest-priority pages.
- 2 **Structure content as Q&A.** Perplexity decomposes user queries into sub-questions and searches for direct answers. Using question-format headings (H2: "How much does X cost?", H3: "What are the key features of X?") aligns perfectly with this behaviour.
- 3 **Load pages with structured data.** Perplexity's parser favours pages with:
 - 4 Comparison tables
 - 5 Numbered lists
 - 6 Clear section headings
 - 7 Statistics with named sources
 - 8 FAQ schema markup
- 9 **Optimise for long-tail, specific queries.** Perplexity users tend to ask highly specific questions. Instead of targeting "best CRM," target "best CRM for 20-person marketing agencies that integrates with HubSpot and Xero."
- 10 **Create "source-worthy" content.** Because Perplexity shows its sources prominently, users can and do click through. Make sure your cited pages deliver on the promise: the content should be comprehensive, current, and visually professional.
- 11 **Robots.txt configuration:**

User-agent: PerplexityBot
Allow: /

Perplexity citation scoring matrix:

Based on analysis of 500+ Perplexity responses across multiple industries, here is what correlates most strongly with citation:

FACTOR	CORRELATION WITH CITATION	ACTIONABLE STEP
Content freshness (updated < 30 days)	Very Strong	Update key pages at least monthly
Presence of statistics/data points	Strong	Include 3+ data points per page
Question-format headings	Strong	Rewrite H2/H3 tags as questions
Domain authority (DR 50+)	Moderate-Strong	Build backlinks through traditional SEO
Table/list formatting	Moderate	Add at least one table per major page
Page load speed (< 2 seconds)	Moderate	Optimise images, use CDN, enable caching
Original research or data	Strong	Publish proprietary surveys, benchmarks, case studies

► 5.4 Google AI Overviews Deep Dive

Understanding Google AI Overviews' source selection:

Google AI Overviews (formerly Search Generative Experience) is unique because it is built on top of Google's existing search infrastructure. This means traditional SEO is a direct prerequisite for AI Overview citation: Google almost exclusively cites pages already ranking in the top 10 organic results.

Google AI Overviews retrieval characteristics:

BEHAVIOUR	DETAIL
Source selection pool	Predominantly from existing top 10 organic results
Citation format	Source cards with site name, favicon, page title, and thumbnail
Query types	Informational, comparative, and local queries most commonly trigger AI Overviews
Schema influence	FAQ, HowTo, and Product schema significantly increase citation probability
Click-through impact	AI Overviews reduce organic CTR by 20-60% for featured queries

Google AI Overviews-specific content tactics:

- 1 **SEO is your entry ticket.** Unlike other AI engines where brand mentions and content quality alone can earn citations, Google AI Overviews requires traditional ranking. If you are not in the top 10 organic results for a query, you will almost never be cited in the AI Overview for that query. Prioritise traditional SEO for your most important keywords.
- 2 **Implement comprehensive schema markup.** Google AI Overviews use schema to understand content structure at a granular level.

Priority schema types for AI Overviews:

SCHEMA TYPE	BEST FOR	AI OVERVIEW IMPACT
FAQPage	Q&A content, knowledge bases	High: directly feeds AI answer generation
HowTo	Tutorials, guides, processes	High: step-by-step content is frequently featured
Product	Product pages, comparisons	High: commercial query answers pull from product schema
Article	Blog posts, news, analysis	Medium: helps identify freshness and authorship

SCHEMA TYPE	BEST FOR	AI OVERVIEW IMPACT
<code>LocalBusiness</code>	Location-based businesses	High: critical for local AI Overview features
<code>Organisation</code>	Company pages, about pages	Medium: establishes entity identity
<code>Review</code> / <code>AggregateRating</code>	Review pages, testimonials	Medium-High: social proof influences recommendation

- 1 **Optimise your Google Business Profile.** For local and commercial queries ("best restaurants near me," "plumber in Manchester"), Google AI Overviews pull heavily from Business Profile data. Ensure yours is complete, current, and has strong reviews.
- 2 **Write concise, direct opening paragraphs.** Google AI Overviews frequently pull the first 1-2 sentences of a page as the core of a cited passage. Make sure your opening paragraph is a standalone, factual answer to the query your page targets.
- 3 **Build authoritative, original content.** Pages with original research, proprietary data, or named expert quotes are cited more frequently than pages that simply aggregate existing information.
- 4 **Robots.txt configuration:**

```
User-agent: Googlebot
Allow: /
```

```
User-agent: Google-Extended
Allow: /
```

Note: `Google-Extended` is the specific user agent for Google's AI training and Gemini. Allowing both `Googlebot` and `Google-Extended` ensures full visibility.

Google AI Overview trigger analysis:

Not all queries generate AI Overviews. Understanding which types do helps you prioritise.

Query Type	AI Overview Frequency	Example	Optimisation Priority
Informational "What is..."	Very High (80%+)	"What is generative engine optimisation?"	Highest: ensure definition-first content
Comparative "Best..."	High (60-70%)	"Best project management tools 2026"	High: build comparison pages
How-to queries	High (60-70%)	"How to improve email deliverability"	High: use HowTo schema, step-by-step format
Local queries	Moderate-High (50-60%)	"Best Italian restaurant in Leeds"	High for local businesses: optimise GBP
Transactional queries	Low-Moderate (20-30%)	"Buy running shoes online"	Lower: Google tends to show shopping results instead
Navigational queries	Low (10-15%)	"Facebook login"	Low: users already know where they are going

6. The 12 Pillars of GEO-Ready Content

These are the actionable strategies that move the needle. Each pillar is grounded in published research, platform testing, and real-world results.

► Pillar 1: Lead with a Direct, Quotable Answer

AI engines scan for concise, extractable answers, typically within the first 100 words of a page. Open every article with a clear, definitional statement.

Rule of thumb: If an AI can't quote your opening paragraph as a standalone answer, rewrite it.

The "Wikipedia Test": Look at how Wikipedia opens its articles. The first sentence is always a clean, factual definition. Apply this principle to every page on your site.

Example for a project management tool:

- "Welcome to ProjectFlow! We're passionate about helping teams do their best work."
- "ProjectFlow is a cloud-based project management platform designed for agencies and consultancies. It combines resource scheduling, time tracking, and client collaboration in a single dashboard, used by over 3,000 agencies across 40 countries."

The second version contains: a clear definition, a target audience, key features, and a credibility signal (3,000 agencies, 40 countries). Any AI engine could extract and cite this as a definitive description.

► **Pillar 2: Structure for Machine Readability**

Use semantic HTML and clear hierarchy:

- **H1** for the page title (one per page)
- **H2** for major sections
- **H3** for sub-topics within sections
- **Bulleted and numbered lists** for steps, features, and comparisons
- **Tables** for structured comparisons
- **Bold text** for key terms and definitions

AI engines parse structured content far more effectively than dense paragraphs.

Content structure scoring guide:

ELEMENT	GEO IMPACT	HOW TO IMPLEMENT
Proper heading hierarchy	High	Every page: H1 > H2 > H3, no skipped levels
Tables	High	At least one data table per major page
Numbered/bulleted lists	Medium-High	Use for any content involving steps, features, or options
Bold key terms	Medium	Bold the first instance of every important term

ELEMENT	GEO IMPACT	HOW TO IMPLEMENT
FAQ sections	High	Add 3-5 Q&A pairs at the bottom of relevant pages
Table of contents	Medium	Add to any page over 1,500 words

► Pillar 3: Embed Statistics and Data Points

The GEO research paper found that content containing statistics and cited data achieved **30-40% higher visibility** in AI-generated responses compared to content without them. Every major claim should be supported by a number.

- INCLUDES "Our product significantly reduces onboarding time."
- INCLUDES "Our product reduces onboarding time by 47%, based on a 2025 study of 1,200 enterprise users (Source: Forrester Research)."

The statistical density benchmark:

CONTENT LENGTH	MINIMUM STATISTICS	OPTIMAL STATISTICS
Under 500 words	2	3-4
500-1,500 words	4	6-8
1,500-3,000 words	8	12-15
Over 3,000 words	12	15-20

Where to find credible statistics:

- Industry analyst reports (Gartner, Forrester, McKinsey, Deloitte)
- Government data (ONS, BLS, Eurostat)
- Academic research papers (Google Scholar)
- Industry association surveys
- Your own customer data and case studies (the most valuable source of all)

► Pillar 4: Include Expert Quotations

Content with attributed expert quotes performs measurably better in generative engines. The original research showed that adding quotations from credible sources improved visibility by up to **25%** in GEO benchmarks.

Name real people. Include their credentials. Make the quotes specific and substantive.

Expert quote formula:

"[Specific, substantive claim]," says **[Full Name]**, [Title] at [Organisation].
"[Supporting context or additional insight]."

Example:

"The biggest mistake brands make in GEO is focusing solely on their own website," says **Dr. Sarah Chen**, Director of Search Research at the Digital Marketing Institute. "AI engines prioritise consensus across multiple sources. If your brand only appears on your own domain, you're invisible to the algorithm."

Finding experts to quote:

SOURCE	HOW TO USE
LinkedIn	Find industry leaders willing to provide quotes for coverage
HARO / Qwoted / Help a B2B Writer	Journalist platforms where experts volunteer quotes
Industry conferences	Quote speakers (with attribution) from public talks
Your own team	Internal subject-matter experts with genuine credentials
Academic researchers	Quote published papers with proper citation

► Pillar 5: Build Topical Authority Through Content Clusters

Do not write one article: build a **content cluster**. A comprehensive hub page linking to 10-20 supporting articles on related sub-topics signals deep expertise to AI engines.

```
Hub: "Complete Guide to Remote Work Tools"
└─ "Best Video Conferencing Platforms 2026"
└─ "How to Choose a Project Management Tool"
└─ "Remote Team Communication: Slack vs Teams vs Discord"
└─ "Asynchronous Work: A Manager's Guide"
└─ "Remote Onboarding: A Step-by-Step Process"
└─ "Cybersecurity for Remote Teams"
└─ "Time Zone Management for Distributed Teams"
└─ ... (10+ more supporting articles)
```

Content cluster scoring:

CLUSTER SIZE	AUTHORITY SIGNAL	TYPICAL AI VISIBILITY LIFT
1 article (standalone)	Minimal	Baseline
3-5 articles (small cluster)	Moderate	10-15% improvement
6-12 articles (medium cluster)	Strong	20-35% improvement
13-20 articles (large cluster)	Very strong	35-50% improvement
20+ articles (comprehensive hub)	Dominant	50%+ improvement

► Pillar 6: Earn Brand Mentions Everywhere

Unlike traditional SEO, **unlinked brand mentions** appear to carry significant weight in GEO. AI engines do not need a hyperlink to register that your brand is associated with a topic; they just need to see your name mentioned in context.

Priority platforms for brand mentions:

PLATFORM	STRATEGY	EFFORT LEVEL	IMPACT
Industry publications	Guest posts, expert commentary, contributed articles	High	Very High
Reddit	Genuine participation in relevant subreddits	Medium	High
YouTube	Video mentions, descriptions, transcripts	Medium	High
Podcast show notes	Guest appearances on relevant podcasts	Medium	Medium-High
Wikipedia	Article creation (if notable) or mention in relevant articles	Very High	Very High
G2 / Trustpilot / Capterra	Encourage authentic customer reviews	Medium	High
Quora / Stack Overflow	Expert answers mentioning your brand naturally	Low-Medium	Medium
LinkedIn articles	Thought leadership from team members	Low	Medium
GitHub	Open-source contributions, documentation	Medium	Medium (for tech brands)

► Pillar 7: Optimise for Conversational Queries

People ask AI engines questions differently than they type into Google. Queries are longer, more conversational, and often multi-part.

Traditional search: "best CRM small business" **AI query:** "I run a 12-person marketing agency in Manchester. What CRM would you recommend and why?"

Structure your content to answer the intent behind the conversation, not just the keyword.

Conversational query patterns to optimise for:

PATTERN	EXAMPLE	CONTENT STRATEGY
Recommendation request	"What would you recommend for..."	Create recommendation guides with criteria-based suggestions
Scenario-based	"I have [specific situation]. What should I..."	Include diverse use cases and scenarios in your content
Comparison	"Which is better, X or Y, for [use case]?"	Build detailed comparison pages with use-case context
Explanation	"Explain [concept] in simple terms"	Write jargon-free explanations alongside technical content
Opinion-seeking	"What do experts think about..."	Include expert quotes and diverse perspectives

► Pillar 8: Publish Original Research and Proprietary Data

Nothing earns AI citations faster than being a **primary source**. If you can publish original surveys, benchmark reports, case studies with real numbers, or industry analysis that no one else has, you become the source that every AI must cite.

Types of original research and their GEO value:

RESEARCH TYPE	EFFORT	GEO VALUE	EXAMPLE
Customer survey	Medium	Very High	"We surveyed 500 UK marketing managers and found that..."
Benchmark report	High	Highest	"2026 Email Marketing Benchmarks: Analysis of 2 Billion Emails"
Case study with metrics	Low-Medium	High	"How [Client] reduced churn by 23% in 90 days"
Industry analysis	Medium	High	"State of AI in UK Financial Services: 2026 Report"

RESEARCH TYPE	EFFORT	GEO VALUE	EXAMPLE
Proprietary data visualisation	Low	Medium-High	Charts and graphs from your own platform data

► Pillar 9: Keep Content Ruthlessly Fresh

Perplexity and Google AI Overviews both heavily favour recent content. Implement a content refresh calendar:

- **Monthly:** Update key statistics, check links, revise any outdated claims
- **Quarterly:** Add new examples, expand sections, incorporate new research
- **Bi-annually:** Restructure and expand based on new queries you are seeing
- **Annually:** Full rewrite with updated research and fresh expert contributions

Include visible "Last updated" dates; both AI crawlers and users value recency signals.

Content freshness priority matrix:

CONTENT TYPE	REFRESH FREQUENCY	WHY
Product comparisons	Monthly	Features and pricing change constantly
Industry statistics	Quarterly	New data releases make old numbers obsolete
How-to guides	Quarterly	Tools and processes evolve
Evergreen definitions	Bi-annually	Core concepts change slowly
Case studies	Annually	Results and context may shift
Technical documentation	As needed	Must match current product/service reality

► Pillar 10: Implement Comprehensive Schema Markup

Structured data helps AI engines understand what **your content is, not just what it says.**

Essential schema types for GEO: - `Article`: publication date, author, headline - `FAQPage`: question-and-answer pairs - `HowTo`: step-by-step processes - `Product`: specifications, pricing, reviews - `Organisation`: brand identity and contact details - `Person`: author expertise and credentials

Schema implementation priority:

SCHEMA TYPE	PRIORITY	PAGES TO APPLY	IMPACT ON GEO
<code>FAQPage</code>	Highest	FAQ pages, knowledge base, product pages with Q&A sections	Direct feed into AI answers
<code>Article</code>	High	All blog posts, guides, and editorial content	Freshness signals, authorship
<code>HowTo</code>	High	Tutorials, process guides, onboarding content	Step-by-step citation
<code>Product</code>	High	Product pages, pricing pages, comparison pages	Commercial query answers
<code>Organisation</code>	Medium	Homepage, about page	Brand entity recognition
<code>Person</code>	Medium	Author pages, team pages	E-E-A-T signals
<code>LocalBusiness</code>	High (if local)	Location pages, contact pages	Local AI Overview features

► Pillar 11: Ensure Technical Accessibility for AI Crawlers

AI crawlers differ from traditional search crawlers. Many struggle with:

- **Client-side JavaScript rendering:** Use server-side rendering (SSR) or static site generation (SSG) where possible
- **Paywalled or gated content:** AI crawlers typically cannot log in or bypass access restrictions

- **Heavily dynamic content:** Ensure critical information is in the initial HTML response

Robots.txt audit: Check that you are not blocking AI crawlers. The major user agents to allow:

User-agent: GPTBot

Allow: /

User-agent: ClaudeBot

Allow: /

User-agent: PerplexityBot

Allow: /

User-agent: Google-Extended

Allow: /

Technical accessibility checklist:

FACTOR	TEST	TOOL
AI crawlers allowed	Check robots.txt for GPTBot, ClaudeBot, PerplexityBot, Google-Extended	Manual review or robots.txt validator
Server-side rendering	View page source (not inspect element); is key content in the HTML?	Browser "View Source"
Page load speed	Under 3 seconds for full page load	Google PageSpeed Insights
Mobile responsive	Content accessible and readable on mobile devices	Google Mobile-Friendly Test
HTTPS	Site uses HTTPS, no mixed content warnings	Browser security indicator
Schema validation	Schema markup is valid and complete	Google Rich Results Test

FACTOR	TEST	TOOL
Sitemap	XML sitemap is current and submitted to Bing Webmaster Tools	Bing Webmaster Tools

► Pillar 12: Write for Synthesis, Not Just Ranking

Traditional SEO content often buries the answer to pad word count. AI engines extract and synthesise; they do not reward length for its own sake.

Write every paragraph as if it might be the only paragraph the AI reads. Make each section self-contained, factual, and valuable in isolation.

The "extraction test": For each major section of your content, ask: "If an AI pulled just this paragraph and showed it to a user, would it make sense on its own? Would it be useful? Would it be accurate?" If the answer to any of these is no, rewrite it.

Synthesis-friendly writing patterns:

PATTERN	EXAMPLE	WHY IT WORKS
Definition-first	"GEO is the practice of..."	AI can extract as a standalone answer
Claim + evidence	"Email marketing ROI averages £36:£1 (DMA, 2025)"	Citable with built-in source
Comparison structure	"Tool A costs £50/mo; Tool B costs £80/mo but includes..."	AI can synthesise comparisons from structured data
Step-by-step	"Step 1: Audit your robots.txt. Step 2: Check schema markup."	Clear sequence for how-to answers

7. Before & After: Real Optimisation Examples

► Example 1: SaaS Product Page

Before (traditional marketing copy):

Welcome to CloudSync, the revolutionary platform that's transforming how teams collaborate in the modern workplace. Our innovative solution leverages cutting-edge technology to deliver an unparalleled experience. With CloudSync, your team will achieve new heights of productivity.

After (GEO-optimised):

CloudSync is a cloud-based team collaboration platform used by over 14,000 organisations across 80 countries. It combines real-time document editing, task management, and video conferencing in a single workspace. In a 2025 benchmark study by Forrester, teams using CloudSync reported a 34% reduction in context-switching and a 28% improvement in project delivery times compared to using separate tools.

Why it works: The optimised version contains a clear definition, specific numbers, a named research source, and factual claims an AI can confidently cite. The original is pure fluff; no AI would ever quote it.

GEO score comparison:

FACTOR	BEFORE	AFTER
Quotable definition	1	1
Specific statistics	1	1 (14,000 orgs, 80 countries, 34%, 28%)
Named source	1	1 (Forrester)
Extractable as standalone	1	1

FACTOR	BEFORE	AFTER
Target audience clarity	■	■ (teams, organisations)

► Example 2: Blog Post Introduction

■ Before:

In today's fast-paced digital landscape, email marketing remains one of the most powerful tools in a marketer's arsenal. But are you really making the most of it? In this comprehensive guide, we'll explore everything you need to know about email marketing in 2026.

■ After:

Email marketing generates an average return of £36 for every £1 spent, making it the highest-ROI channel in digital marketing (DMA, 2025). This guide covers the five strategies that drove the best results in 2025-26: segmentation, behavioural triggers, interactive emails, AI-powered personalisation, and deliverability optimisation. Each section includes benchmarks, examples, and implementation steps.

Why it works: The optimised version opens with a cited statistic, previews specific content, and reads as an authoritative reference. An AI answering "What is the ROI of email marketing?" would cite this immediately.

► Example 3: FAQ Section

■ Before:

Q: How long does shipping take? A: Shipping times vary depending on your location and the shipping method you choose. We offer several options to meet your needs.

After:

Q: How long does shipping take? A: Standard UK delivery takes 3-5 working days. Express delivery arrives within 1-2 working days. International shipping to Europe takes 5-10 working days and worldwide 10-21 working days. All orders include tracking. Orders placed before 2:00 PM GMT on weekdays are despatched the same day.

Why it works: Specificity. AI engines prioritise content that answers questions completely **in** a single passage. Vague answers are never cited.

► **Example 4: Professional Services "About" Page**

Before:

At Smith & Partners, we pride ourselves on delivering exceptional legal services. Our team of dedicated professionals is committed to achieving the best outcomes for our clients. With decades of combined experience, we bring a wealth of knowledge to every case.

After:

Smith & Partners is a commercial law firm based in Manchester, specialising in technology transactions, data protection (UK GDPR), and intellectual property licensing. Founded in 2008, the firm has 34 solicitors across three offices (Manchester, London, Edinburgh) and advises over 200 technology companies annually, from seed-stage startups to FTSE 250 enterprises. The firm was ranked in Legal 500 (2025) for Technology and Telecoms and holds a Chambers UK Band 2 ranking for IP in the North West.

Why it works: The original says nothing an AI could cite. The revised version is packed with extractable facts: location, specialisms, founding year, team size, client volume, and named third-party rankings. When an AI is asked "Which law firms in Manchester specialise in technology law?", this version provides everything it needs to include Smith & Partners in the answer.

GEO score comparison:

FACTOR	BEFORE	AFTER
Location specified	■	■ (Manchester, London, Edinburgh)
Specialisms named	■	■ (tech transactions, GDPR, IP)
Verifiable credentials	■	■ (Legal 500, Chambers UK)
Quantified scale	■	■ (34 solicitors, 200+ clients)
Founded date	■	■ (2008)

► Example 5: E-commerce Product Description

■ **Before:**

Our amazing noise-cancelling headphones deliver an incredible listening experience. With state-of-the-art technology and premium materials, these headphones are perfect for music lovers who demand the best. Available in three stylish colours.

After:

The **SonicPro NC-700** noise-cancelling headphones use hybrid active noise cancellation with four microphones per ear cup, reducing ambient noise by up to 38dB. Battery life is rated at 32 hours with ANC enabled and 45 hours without. They weigh 254g, support Bluetooth 5.3 with multipoint connection (two devices simultaneously), and include USB-C fast charging (15 minutes for 5 hours of playback). Rated 4.6/5 across 2,300 reviews on Trustpilot. Available in Graphite Black, Arctic White, and Navy Blue. RRP: £249.

Why it works: The original is undifferentiated marketing copy. The revised version is a specification sheet that AI engines can parse, compare, and cite. When a user asks "What are the best noise-cancelling headphones under £300?", the AI has concrete data to work with: specific dB reduction, battery life, weight, Bluetooth version, and price. These are the exact data points that AI engines extract for comparison answers.

GEO score comparison:

FACTOR	BEFORE	AFTER
Product model name	■	■ (SonicPro NC-700)
Technical specifications	■	■ (38dB, 32hrs, 254g, BT 5.3)
Social proof	■	■ (4.6/5, 2,300 reviews)
Price	■	■ (£249)
Comparable data points	■	■ (every spec is comparable)

► Example 6: B2B Case Study

Before:

"Working with DataPulse transformed our business. Their team was incredibly professional and the results speak for themselves. We couldn't be happier with the partnership." , Happy Client

After:

Case Study: How Meridian Logistics Reduced Freight Costs by 19% with DataPulse

Client: Meridian Logistics, a UK-based 3PL provider with £45M annual revenue and 120 employees. **Challenge:** Manual route planning was causing an estimated £780,000 in unnecessary fuel costs annually. **Solution:** DataPulse's AI-powered route optimisation platform, deployed across Meridian's fleet of 340 vehicles over a 12-week rollout. **Results:** 19% reduction in fuel costs (£148,000 annual saving), 23% improvement in delivery punctuality, and ROI achieved in 4.5 months. "DataPulse paid for itself before the end of Q2. The route optimisation alone saved us more than the entire annual licence fee," says **James Whitfield**, Operations Director at Meridian Logistics.

Why it works: The original is a generic testimonial that no AI would cite. The revised version is a structured case study with named client, specific challenge, measurable results, and an attributed quote. When an AI is asked "What are the benefits of AI route optimisation for logistics companies?", this case study provides exactly the kind of evidence it needs.

8. The Content Audit Framework

This framework provides a systematic, step-by-step process for auditing your existing content for GEO readiness. Use it to evaluate and prioritise improvements across your entire site.

► Step 1: Identify Your Priority Pages

Not every page needs GEO optimisation. Start with the pages that drive the most business value.

Priority page selection criteria:

CRITERION	HOW TO ASSESS	WEIGHT
Revenue impact	Pages that directly drive leads, sales, or conversions	★★★★★
Search volume	Pages targeting keywords with significant monthly search volume	★★★★☆
Current ranking	Pages already ranking in top 20 (easier to optimise than starting from scratch)	★★★★☆
AI query relevance	Pages targeting topics that users commonly ask AI engines about	★★★★☆
Competitive gap	Topics where competitors are being cited but you are not	★★★★★

Action: List your top 20 pages by business impact. These are your GEO audit candidates.

► Step 2: Run the GEO Readiness Score

For each priority page, score it against the following criteria. Each item scores 0 (absent), 1 (partial), or 2 (fully present).

GEO Readiness Scorecard:

#	CRITERION	0 (ABSENT)	1 (PARTIAL)	2 (FULLY PRESENT)	SCORE
1	Quotable opening paragraph	No clear definition or direct answer	Definition exists but lacks specificity	Clean, extractable definition with key facts	/2
2	Statistics and data points	No statistics	1-2 stats without named sources	3+ stats with named, credible sources	/2

#	CRITERION	0 (ABSENT)	1 (PARTIAL)	2 (FULLY PRESENT)	SCORE
3	Expert quotes	No expert quotes	Generic testimonials	Named experts with credentials and specific claims	/2
4	Heading structure	Poor or no heading hierarchy	Some headings but inconsistent	Full H1 > H2 > H3 hierarchy, no skipped levels	/2
5	Tables and lists	No tables or lists	Some formatting	Tables for comparisons, lists for features/steps	/2
6	FAQ section	No FAQ	Partial FAQ	Comprehensive FAQ with question-format headings	/2
7	Schema markup	No schema	Basic schema (Article only)	Comprehensive schema (FAQ, HowTo, Product, etc.)	/2
8	Freshness signals	No update date, content over 12 months old	Updated within 12 months	Updated within 90 days with visible date	/2
9	External brand mentions	Fewer than 3 external mentions for the target topic	3-10 external mentions	10+ external mentions across diverse sources	/2
10	AI crawler access	AI crawlers blocked in robots.txt	Some crawlers allowed	All major AI crawlers explicitly allowed	/2
11	Self-contained sections	Sections require context from other parts of the page	Most sections standalone	Every section extractable as a standalone answer	/2

#	CRITERION	0 (ABSENT)	1 (PARTIAL)	2 (FULLY PRESENT)	SCORE
12	Conversational query match	Content targets keywords only	Some conversational headings	Content structured around natural language questions	/2

Scoring interpretation:

SCORE	RATING	ACTION
20-24	Excellent	Minor tweaks only; focus on external authority building
15-19	Good	Targeted improvements to specific weak areas
10-14	Needs Work	Significant restructuring and content additions required
5-9	Poor	Major overhaul required; consider rewriting from scratch
0-4	Not GEO-Ready	Full rewrite needed with GEO principles from the ground up

► Step 3: Conduct the Technical Audit

For each priority page, check the following technical factors:

Technical audit walkthrough:

- 1 **Check robots.txt (5 minutes)**
- 2 Visit yourdomain.com/robots.txt
- 3 Search for: GPTBot, ClaudeBot, PerplexityBot, Google-Extended
- 4 If any are blocked (Disallow), update immediately
- 5 If not mentioned, they are allowed by default (good), but explicitly allowing them is better practice
- 6 **Validate schema markup (10 minutes per page)**
- 7 Use Google's Rich Results Test: search.google.com/test/rich-results
- 8 Paste your page URL

- 9 Check for: valid FAQ, Article, HowTo, Product schema as appropriate
- 10 Note any errors or warnings for fixing
- 11 **Test rendering (5 minutes per page)**
- 12 Right-click > View Page Source (not Inspect Element)
- 13 Search for your key content in the raw HTML
- 14 If your main content is not visible in the source, it is JavaScript-rendered and potentially invisible to AI crawlers
- 15 Solution: implement server-side rendering or pre-rendering
- 16 **Verify page speed (5 minutes per page)**
- 17 Run through Google PageSpeed Insights
- 18 Target: LCP under 2.5 seconds, CLS under 0.1, INP under 200ms
- 19 Fix critical speed issues that could prevent AI crawlers from fully loading the page
- 20 **Check mobile responsiveness (5 minutes per page)**
- 21 Test on Google's Mobile-Friendly Test
- 22 Ensure all content, tables, and media display correctly on mobile

► Step 4: Audit Content Quality

Walk through each priority page with these specific checks:

Opening paragraph audit:

- [] Does the first sentence define what the page is about?
- [] Could the opening paragraph be quoted as a standalone answer?
- [] Does it contain at least one specific fact or statistic?
- [] Does it name the brand and category clearly?

Statistical density audit:

- [] Count the total number of statistics on the page
- [] For each statistic, is the source named?
- [] Are statistics from the current or previous year?
- [] Do statistics support the page's primary claims?

Structure audit:

- [] Map the heading hierarchy (H1 > H2 > H3): are there gaps?
- [] Count tables: is there at least one per major section?
- [] Count lists: are features, steps, and options formatted as lists?
- [] Is there an FAQ section?

Expert content audit:

- [] Are there expert quotes with names and credentials?
- [] Does the page demonstrate first-hand experience?
- [] Is the author identified with a bio and credentials?

► Step 5: Audit External Authority

For each priority topic (not page, but topic), assess your external brand presence:

Brand mention audit process:

- 1 Search for `"your brand name" + "topic keyword"` across:
- 2 Google (regular search)
- 3 Reddit (via Google: `site:reddit.com "your brand" "topic"`)
- 4 YouTube (search for brand mentions in video titles and descriptions)
- 5 Major industry publications in your sector
- 6 Count the total number of unique external sources mentioning your brand in relation to the topic
- 7 Identify gaps: which platforms and publications mention competitors but not you?

External authority scorecard:

SOURCE TYPE	YOUR MENTIONS	TOP COMPETITOR MENTIONS	GAP
Industry publications	?	?	?
Reddit threads	?	?	?

Source Type	Your Mentions	Top Competitor Mentions	GAP
YouTube mentions	?	?	?
Review platforms (G2, etc.)	?	?	?
Wikipedia	?	?	?
Podcast appearances	?	?	?
Total	?	?	?

► Step 6: Prioritise and Plan

With scores complete, prioritise your pages using this matrix:

PRIORITY	CRITERIA	ACTION TIMELINE
P1: Quick Wins	GEO Score 15-19, already ranking well, small gaps to fill	Week 1-2: optimise these first
P2: High Impact	GEO Score 10-14, high business value, requires significant work	Week 2-4: restructure and enhance
P3: Rebuilds	GEO Score under 10, high business value, needs complete overhaul	Month 2: full rewrite
P4: New Content	Topic gaps where you have no content but competitors are being cited	Month 2-3: create new GEO-ready content
P5: Low Priority	Low business value or low AI query relevance	Backlog: address when higher priorities are complete

9. The GEO Audit Checklist

Use this checklist to audit any page on your site for GEO readiness. This is the condensed, reference version of the full Content Audit Framework in Section 8.

► Content Quality

- [] Opening paragraph contains a clear, quotable definition or answer
- [] Key claims are supported with statistics, percentages, or named data sources
- [] At least 3 statistics with named sources per 1,000 words
- [] Expert quotes with attributed names and credentials are included
- [] Content answers the query directly: no filler before the answer
- [] Each section is self-contained and extractable
- [] Content is written in British English consistently (not mixed)
- [] "Last updated" date is visible and accurate
- [] Content includes at least one original insight, data point, or framework not available elsewhere

► Content Structure

- [] Proper heading hierarchy (H1 > H2 > H3, no skipped levels)
- [] Bulleted or numbered lists used for steps, features, and comparisons
- [] At least one data table per major page
- [] Key terms and definitions are bold or highlighted
- [] FAQ section included with question-format headings
- [] Table of contents included for pages over 1,500 words
- [] Content sections could be extracted independently and still make sense

► Technical Foundation

- [] AI crawlers (GPTBot, ClaudeBot, PerplexityBot, Google-Extended) are allowed in robots.txt
- [] Page uses server-side rendering or pre-rendering
- [] Schema markup implemented (Article, FAQPage, HowTo, Product, Organisation as relevant)
- [] Page loads without requiring JavaScript for critical content
- [] Core Web Vitals passing (LCP < 2.5s, CLS < 0.1, INP < 200ms)
- [] HTTPS enabled

- [] Mobile responsive
- [] XML sitemap is current and submitted to Bing Webmaster Tools
- [] No broken internal or external links

► Authority and Distribution

- [] Brand is mentioned on at least 10 external authoritative sources for the target topic
- [] Content is referenced or linked from industry publications
- [] Brand has presence on UGC platforms (Reddit, Quora, YouTube, G2)
- [] Google Business Profile is complete and optimised (for local/commercial queries)
- [] Wikipedia article exists or brand is mentioned in relevant Wikipedia articles
- [] Social proof (reviews, testimonials, case studies) is publicly accessible
- [] Brand appears in at least 3 different source trust tiers (see Section 3.4)

► Freshness

- [] Content has been updated within the last 90 days
 - [] Statistics and data points reference the current or previous year
 - [] No broken links or references to discontinued products/services
 - [] Content refresh schedule is documented and followed
 - [] "Last updated" date reflects the most recent substantive edit
-

10. Implementation Roadmap: Your First 30 Days

This section provides a week-by-week action plan for implementing GEO across your organisation. Each week builds on the previous one.

► Week 1: Foundation and Assessment

Goal: Understand your current AI visibility and establish your baseline.

DAY	TASK	TIME ESTIMATE	DELIVERABLE
Mon	Set up AI visibility tracking: create a spreadsheet with 20-30 target prompts across all four engines	2-3 hours	Tracking spreadsheet with baseline data
Mon	Run each prompt across ChatGPT, Claude, Perplexity, and Google; record brand mentions, citations, and competitor presence	3-4 hours	Completed baseline assessment
Tue	Audit robots.txt for all AI crawlers; fix any blocks immediately	30 minutes	Updated robots.txt
Tue	Identify your top 20 pages by business value (revenue impact, traffic, strategic importance)	1-2 hours	Prioritised page list
Wed-Thu	Run the GEO Readiness Scorecard (Section 8, Step 2) on all 20 priority pages	3-4 hours	Scored pages with identified gaps
Fri	Conduct the brand mention audit (Section 8, Step 5) for your top 5 topics	2-3 hours	External authority scorecard
Fri	Compile findings into a GEO Action Plan document	1-2 hours	Action plan with prioritised tasks

Week 1 deliverables: - AI visibility baseline data - Robots.txt fixed (if needed) - 20 pages scored for GEO readiness - External authority audit for top 5 topics - Prioritised action plan for Weeks 2-4

► Week 2: Quick Wins and Technical Fixes

Goal: Implement the changes with the highest impact-to-effort ratio.

DAY	TASK	TIME ESTIMATE	DELIVERABLE
Mon	Rewrite opening paragraphs of your top 10 pages to lead with quotable definitions	3-4 hours	10 updated opening paragraphs
Tue	Add statistics with named sources to all 10 pages (minimum 3 per page)	3-4 hours	30+ statistics added
Wed	Implement FAQ schema on your top 5 pages; add FAQ sections where missing	3-4 hours	FAQ schema live on 5 pages
Thu	Add at least one data table to each of your top 10 pages	2-3 hours	10 new data tables
Thu	Implement Article schema on all blog posts that lack it	2-3 hours	Article schema deployed
Fri	Add visible "Last updated" dates to all priority pages; update any statistics older than 12 months	2-3 hours	Freshness signals added

Week 2 deliverables: - 10 pages with GEO-optimised opening paragraphs - 30+ statistics added across priority pages - FAQ schema on 5 pages - 10 new data tables - Freshness signals on all priority pages

► Week 3: Content Depth and Authority Building

Goal: Build the content depth and external authority that drive AI citations.

DAY	TASK	TIME ESTIMATE	DELIVERABLE
Mon	Identify 3 expert sources for quotes; reach out via LinkedIn, HARO, or your network	2-3 hours	Expert outreach initiated
Mon-Tue	Write 2 new supporting articles for your highest-priority content cluster	6-8 hours	2 new cluster articles published

DAY	TASK	TIME ESTIMATE	DELIVERABLE
Wed	Submit guest post pitches to 5 industry publications in your sector	2-3 hours	5 pitches sent
Wed	Create or update your Wikipedia presence (if brand is notable) or identify relevant Wikipedia articles to contribute to	2-3 hours	Wikipedia strategy documented
Thu	Begin Reddit participation strategy: identify 3-5 relevant subreddits, contribute genuinely helpful answers that naturally mention your brand	2-3 hours	Reddit engagement started
Fri	Reach out to 3 relevant podcasts for guest appearances	1-2 hours	Podcast pitches sent
Fri	Encourage 10 existing customers to leave reviews on G2, Trustpilot, or relevant review platforms	1-2 hours	Review request emails sent

Week 3 deliverables: - Expert outreach in progress - 2 new cluster articles published - 5 guest post pitches sent - Reddit engagement started - Podcast outreach initiated - Customer review requests sent

► Week 4: Measurement, Iteration, and Systemisation

Goal: Measure progress, iterate on what is working, and build sustainable systems.

DAY	TASK	TIME ESTIMATE	DELIVERABLE
Mon	Re-run all 20-30 AI visibility prompts; compare to Week 1 baseline	3-4 hours	Week 4 visibility data
Mon	Calculate improvement: how many new citations, mentions, and visibility gains?	1-2 hours	Progress report

DAY	TASK	TIME ESTIMATE	DELIVERABLE
Tue	Based on results, identify which optimisations had the most impact; double down on those	2-3 hours	Updated strategy priorities
Wed	Optimise 5 more pages using the most effective tactics from Weeks 2-3	4-5 hours	5 additional pages optimised
Thu	Set up a monthly content refresh calendar for all priority pages	1-2 hours	Content refresh calendar
Thu	Document your GEO process: create internal guidelines so any team member can optimise new content	2-3 hours	Internal GEO guidelines document
Fri	Plan Month 2: identify next batch of pages, new content to create, and authority-building actions	2-3 hours	Month 2 action plan

Week 4 deliverables: - Week 4 visibility data (compared to baseline) - Progress report with measurable improvements - 5 additional pages optimised - Monthly content refresh calendar - Internal GEO guidelines - Month 2 plan

► 30-Day Milestone Targets

METRIC	WEEK 1 BASELINE (TYPICAL)	WEEK 4 TARGET	HOW TO MEASURE
Pages GEO-optimised	0	15-20	Count of pages that score 15+ on GEO Readiness Scorecard
AI citation rate	5-15% of target prompts	20-30% of target prompts	Manual testing across all four engines

METRIC	WEEK 1 BASELINE (TYPICAL)	WEEK 4 TARGET	HOW TO MEASURE
External brand mentions (for top topic)	Varies	+10 new mentions	Brand mention audit (Section 8, Step 5)
Schema markup coverage	Varies	100% of priority pages	Rich Results Test
Content freshness	Varies	100% of priority pages updated within 30 days	Visible "Last updated" dates

11. Measuring AI Visibility

You cannot improve what you cannot measure. Here is how to track your GEO performance systematically.

► Manual Monitoring

The prompt test: Regularly ask each AI engine questions that your brand should appear in.

Track:

- Whether your brand is mentioned
- Whether your website is cited as a source
- What position your brand appears in (first mentioned? third?)
- What competitors are cited instead
- How the AI describes your brand (accurate? favourable? outdated?)

Create a tracking spreadsheet:

PROMPT	ENGINE	BRAND MENTIONED?	POSITION	COMPETITORS CITED	SENTIMENT	DATE
"Best CRM for agencies UK"	ChatGPT	Yes	2nd	HubSpot, Pipedrive	Positive	03/02/2024
"Best CRM for agencies UK"	Perplexity	No	N/A	Salesforce, Zoho	N/A	03/02/2024
"Best CRM for agencies UK"	Claude	Yes	3rd	HubSpot, Salesforce	Neutral	03/02/2024
"Best CRM for agencies UK"	Google AIO	Yes	1st	HubSpot	Positive	03/02/2024

Prompt design best practices:

PROMPT TYPE	PURPOSE	EXAMPLE
Category prompt	Test if brand appears for the broad category	"What are the best [category] tools?"
Feature prompt	Test for specific capability association	"Which [category] tool is best for [feature]?"
Comparison prompt	Test head-to-head positioning	"Compare [your brand] vs [competitor]"
Recommendation prompt	Test scenario-based citation	"I need [specific use case]. What would you recommend?"
Reputation prompt	Test brand perception accuracy	"Tell me about [your brand]"
Local prompt	Test geographic association	"Best [category] in [location]"

► Automated Tools

Several platforms now offer AI visibility tracking. Here is a comparison of the leading options:

TOOL	WHAT IT TRACKS	PRICE RANGE (MONTHLY)	BEST FOR
Semrush AI Visibility Toolkit	Brand mentions, sentiment, share of voice across AI engines	£100-300 (part of Semrush subscription)	Comprehensive tracking alongside traditional SEO
Profound	Brand presence in LLM responses across GPT, Claude, Gemini	£150-500	Dedicated AI visibility monitoring
Scrunch AI	AI citation patterns, competitor analysis, recommendation tracking	£100-400	Competitor intelligence
Otterly.AI	Brand appearances in AI-generated answers with historical tracking	£80-250	Affordable entry-level tracking
Peec AI	AI search results monitoring with alert system	£50-200	Budget-friendly, alert-focused
Manual tracking (spreadsheet)	Whatever you choose to test	Free (time cost only)	Small businesses, early-stage GEO

► Key Metrics to Track

METRIC	DEFINITION	TARGET	MEASUREMENT FREQUENCY
AI Share of Voice	% of relevant prompts where your brand is cited vs. competitors	> 30%	Monthly
Citation Frequency	How often your domain appears as a cited source	Increasing month-on-month	Monthly

METRIC	DEFINITION	TARGET	MEASUREMENT FREQUENCY
Brand Sentiment	Whether AI descriptions of your brand are positive, neutral, or negative	Positive	Monthly
Source Accuracy	Whether AI engines present accurate information about your brand	> 95%	Monthly
Prompt Coverage	% of target prompts where your brand appears at all	> 50%	Monthly
Position in Response	Where your brand appears (1st, 2nd, 3rd mentioned)	Top 3	Monthly
Cross-Engine Consistency	Whether your brand appears consistently across all four engines	All four	Monthly

► Setting Up Your Measurement Dashboard

Recommended dashboard structure:

SECTION	METRICS	DATA SOURCE
Overview	Overall AI citation rate, month-on-month trend, top engine performance	Aggregated from all tracking
Engine Breakdown	Citation rate per engine (ChatGPT, Claude, Perplexity, Google AIO)	Per-engine testing
Competitor Comparison	Your share of voice vs. top 3 competitors	Competitive prompt testing
Content Performance	Which pages are being cited most frequently	Source URL tracking
Accuracy Tracker	Any inaccurate AI descriptions of your brand	Reputation prompts

SECTION	METRICS	DATA SOURCE
Action Items	Specific optimisation tasks based on measurement gaps	Analysis of weak points

12. GEO vs. SEO: Allies, Not Rivals

A common misconception is that GEO replaces SEO. It does not. They are deeply complementary, and the most effective strategy treats them as two aspects of the same programme.

► What SEO Does for GEO

- **Ranking in Google's top 10 is a prerequisite** for being cited in AI Overviews: if you do not rank, you will not be cited
- **Backlinks build the domain authority** that AI engines use as a trust signal during source retrieval
- **Keyword research identifies the queries** that inform your GEO content strategy
- **Technical SEO ensures crawlability** for both traditional and AI crawlers

► What GEO Adds Beyond SEO

- **Unlinked mention strategy:** building brand presence where links do not matter
- **Cross-platform optimisation:** ensuring visibility in ChatGPT, Claude, and Perplexity, not just Google
- **Synthesis-first writing:** structuring content for extraction rather than just ranking
- **AI crawler management:** specifically managing access for GPTBot, ClaudeBot, and PerplexityBot
- **Prompt-based strategy:** optimising for conversational queries, not just keyword strings

► SEO vs. GEO Comparison Matrix

DIMENSION	TRADITIONAL SEO	GEO	OVERLAP
Primary goal	Rank in organic results	Get cited in AI answers	Both increase visibility
Key metric	Ranking position, organic traffic	AI citation rate, share of voice	Both drive brand awareness and leads
Content format	Long-form, keyword-optimised	Structured, extractable, fact-dense	Both reward quality content
Link building	Hyperlinks from authoritative sites	Unlinked brand mentions across the web	Links help both; mentions help GEO more
Technical	Crawlability, Core Web Vitals, mobile	AI crawler access, schema, SSR	Technical foundations serve both
Freshness	Important but not critical for all queries	Critical, especially for Perplexity and Google AIO	Both benefit from regular updates
Target keyword	Short-tail and long-tail keywords	Conversational queries and prompts	Long-tail SEO aligns well with GEO
Measurement	Google Search Console, rank trackers	AI prompt testing, citation monitoring	Some tools now cover both

► The Unified Approach

The most effective strategy is to treat GEO as an **extension** of your existing SEO programme. Every SEO improvement you make, from better content to more backlinks to a stronger technical foundation, also improves your GEO performance. The GEO-specific additions (structured answers, statistical density, expert quotes, unlinked mentions, AI crawler access) layer on top.

The integration model:

- Traditional SEO Foundation
- └ Keyword research → Informs GEO prompt strategy
 - └ Content creation → Apply GEO pillars to all new content
 - └ Link building → Add unlinked mention strategy for GEO
 - └ Technical SEO → Extend to AI crawler management
 - └ Analytics → Add AI visibility monitoring
 - └ Content refresh → Align with GEO freshness requirements

► Budget Allocation Guidance

For organisations wondering how to split resources between SEO and GEO:

ORGANISATION TYPE	RECOMMENDED SEO : GEO SPLIT	REASONING
Early-stage (weak SEO foundation)	80% SEO : 20% GEO	Build the foundation first; SEO directly supports GEO
Established (strong SEO, weak AI visibility)	50% SEO : 50% GEO	SEO maintenance plus aggressive GEO build-out
Mature (strong SEO, some AI visibility)	40% SEO : 60% GEO	GEO is where the marginal gains are largest
Advanced (strong SEO and GEO)	30% SEO : 70% GEO	SEO in maintenance mode; GEO is the growth frontier

13. What's Coming Next

GEO is in its infancy. Here is what we expect to see in the next 12-24 months, and how to prepare now.

► 13.1 AI Engine Advertising

Paid placements within AI-generated responses are coming. Perplexity has already begun testing sponsored results. Google will follow. Microsoft has hinted at ad formats within Copilot.

What this means for you: Organic GEO (earning citations through content quality and authority) will become even more valuable as paid options emerge. Just as organic SEO became more valuable when Google Ads made the top of the SERP paid-for, organic AI citations will carry a premium of trust.

How to prepare: Build your organic AI visibility now, before advertising makes the space more competitive and expensive.

► 13.2 Standardised AI Analytics

Just as Google Analytics became the standard for web traffic, tools for measuring AI citation performance will mature rapidly. Expect:

- Dedicated AI visibility dashboards integrated into existing SEO platforms
- Standardised metrics for AI share of voice
- Attribution models that connect AI citations to conversions

How to prepare: Start measuring now, even with manual methods. The habit of tracking AI visibility will put you ahead when automated tools mature.

► 13.3 AI-Specific Content Formats

We expect new structured data standards designed specifically for AI consumption, beyond current schema.org markup. These might include:

- Machine-readable "answer blocks" that AI engines can extract directly
- Standardised citation metadata (how your content should be attributed)
- AI-optimised sitemaps with content summaries

How to prepare: Implement comprehensive schema now. Organisations that are already structured-data-mature will adopt new standards faster.

► 13.4 Personalised AI Results

As AI engines learn individual user preferences, brand visibility will become more nuanced. Your content may be shown to some users and hidden from others based on their history.

How to prepare: Build broad authority across multiple topics and platforms. The more diverse your presence, the more user segments you will reach.

► 13.5 Voice and Multimodal Search

As AI assistants become the primary interface through smart speakers, in-car systems, and AR glasses, GEO will expand to optimise for voice synthesis and visual responses.

Key difference for voice: Voice answers are even more binary than text. A smart speaker gives one answer, not a list. Being the single cited source becomes even more critical.

How to prepare: Write content that sounds natural when read aloud. Test your key definitions by reading them out loud: if they sound stilted, rewrite them.

► 13.6 Regulatory Scrutiny

Governments are beginning to ask whether AI engines have an obligation to fairly represent all sources, not just the ones they prefer. The EU AI Act, UK Online Safety Act, and similar legislation may reshape how citations work.

How to prepare: Focus on genuine quality and authority. Regulatory changes will likely penalise manipulation and reward authentic expertise, which is exactly what good GEO already emphasises.

► 13.7 Timeline of Expected Changes

DEVELOPMENT	EXPECTED TIMELINE	IMPACT LEVEL	PREPARATION PRIORITY
AI advertising (Perplexity, Google)	Already started, scaling in 2026	High	Build organic visibility now
Standardised AI analytics	H2 2026 - H1 2027	Medium	Start manual tracking immediately

Development	Expected Timeline	Impact Level	Preparation Priority
AI-specific content formats	2027	Medium	Master current schema first
Personalised AI results	2026-2027 (gradual)	High	Build diverse authority
Voice-first AI search	2026-2028	High for B2C	Optimise for spoken delivery
Regulatory frameworks	2027-2028	Uncertain	Focus on genuine quality

Final Thoughts

The transition from traditional search to AI-powered discovery is not a distant possibility; it is happening now. Every month, more users turn to ChatGPT, Claude, Perplexity, and Google AI Overviews for the answers that used to require clicking through ten blue links.

The brands that win in this new landscape will be those that understand a simple truth: **AI engines don't rank pages. They cite sources.** And to be cited, your content must be authoritative, structured, specific, evidence-based, and present across the web.

GEO is not a replacement for good marketing. It is the evolution of it. The fundamentals, including genuine expertise, valuable content, and a trustworthy reputation, matter more than ever. GEO simply gives you the framework to ensure those fundamentals are visible to the machines that are increasingly deciding which brands get recommended and which get ignored.

Your action plan starts now:

- 1 Run the baseline assessment (Week 1 of the Implementation Roadmap)
- 2 Fix the technical foundations (robots.txt, schema, crawlability)
- 3 Optimise your top 10 pages using the 12 Pillars
- 4 Build external authority through mentions, reviews, and expert content
- 5 Measure monthly and iterate

The first-mover advantage in GEO is real, and the window is open. Every week you wait, your competitors may be the ones getting cited instead.

14. Appendix: Resources, Tools, and Further Reading

► Essential Reading

RESOURCE	TYPE	WHAT YOU WILL LEARN
"GEO: Generative Engine Optimization" (Aggarwal et al., 2023)	Research paper	The foundational academic research behind GEO; published at KDD 2024
Google Search Central Blog	Official documentation	How Google AI Overviews select and cite sources
Bing Webmaster Guidelines	Official documentation	How Bing (ChatGPT's search backbone) ranks content
Perplexity AI Blog	Official documentation	How Perplexity indexes and cites the web
Schema.org Documentation	Technical standard	Full reference for structured data markup

► Recommended Tools

TOOL	PURPOSE	FREE/PAID
Google Search Console	Monitor organic rankings and indexing	Free
Bing Webmaster Tools	Monitor Bing indexing (critical for ChatGPT)	Free
Google Rich Results Test	Validate schema markup	Free
Google PageSpeed Insights	Test page load performance	Free

TOOL	PURPOSE	FREE/PAID
Schema Markup Generator (Merkle)	Generate schema code	Free
Semrush / Ahrefs	Comprehensive SEO and emerging GEO metrics	Paid
Screaming Frog	Technical site audit	Free (limited) / Paid

► AI Crawler User Agents Reference

CRAWLER	USER AGENT STRING	OWNER	PURPOSE
GPTBot	<code>GPTBot</code>	OpenAI	Training data and web browsing for ChatGPT
ChatGPT-User	<code>ChatGPT-User</code>	OpenAI	Live web browsing during ChatGPT conversations
ClaudeBot	<code>ClaudeBot</code>	Anthropic	Training data and web access for Claude
PerplexityBot	<code>PerplexityBot</code>	Perplexity	Real-time web indexing for Perplexity answers
Googlebot	<code>Googlebot</code>	Google	Standard web crawling for Google Search and AI Overviews
Google-Extended	<code>Google-Extended</code>	Google	AI training data specifically for Gemini and AI features

► Recommended Robots.txt Configuration

```
# Allow all AI crawlers for maximum GEO visibility
User-agent: GPTBot
Allow: /

User-agent: ChatGPT-User
Allow: /

User-agent: ClaudeBot
Allow: /

User-agent: PerplexityBot
Allow: /

User-agent: Google-Extended
Allow: /

# Standard Googlebot (already allowed by default in most configs)
User-agent: Googlebot
Allow: /
```

End of Playbook **The GEO Playbook** · v2.0 · February 2026 *© 2026. All rights reserved.*